

Koneru Lakshmaiah Education Foundation
(Deemed to be University)

A Project Based Lab Report
On
Monitoring and Logging System
with Amazon CloudWatch

BACHELOR OF TECHNOLOGY

SUBMITTED BY

Bodepudi Jothirmaye

2200032386



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

KONERU LAKSHMAIAH EDUCATION FOUNDATION

Green Fields, Vaddeswaram,
Guntur, Andhra Pradesh, 522 302
INDIA

ACKNOWLEDGEMENTS

It is great pleasure for me to express my gratitude to our honorable President **Sri. Koneru Satyanarayana**, for giving the opportunity and platform with facilities in accomplishing the project based laboratory report.

I express the sincere gratitude to our principal **Dr. A. Jagadeesh** for his administration towards our academic growth.

I express sincere gratitude to HOD-CSE **Prof. VSV Prabhakar** for his leadership and constant motivation provided in successful completion of our academic semester. I record it as my privilege to deeply thank for providing us the efficient faculty and facilities to make our ideas into reality.

I express my sincere thanks to our project supervisor Mrs. Sunitha Bandeela for her novel association of ideas, encouragement, appreciation and intellectual zeal which motivated us to venture this project successfully.

Finally, it is pleased to acknowledge the indebtedness to all those who devoted themselves directly or indirectly to make this project report success.

ABSTRACT

Serverless image processing with AWS Lambda and S3 offers a highly scalable, cost-effective, and automated approach to handling image transformations in the cloud. By utilizing AWS Lambda, images can be processed dynamically in response to S3 events, eliminating the need for server management. Integration with additional AWS services, such as Amazon SNS and Amazon CloudWatch, enhances workflow automation, monitoring, and alerting.

This method is particularly beneficial for real-time applications such as thumbnail generation, watermarking, and format conversion. It optimizes infrastructure usage, reduces operational costs, and improves scalability. However, considerations such as Lambda execution time, cold start mitigation, and secure access management are crucial for efficient implementation.

Overall, serverless image processing enables businesses to achieve high availability, fault tolerance, and enhanced performance, making it an ideal choice for modern cloud-based applications.

INDEX

S NO.	NAME	PAGE NO.
1.	Project Aim	5
2.	Description	6-7
3.	Implementation steps	8-25
4.	Output	26-27
5.	Conclusion	28

PROJECT AIM

In this project, you will create a serverless image processing pipeline that automatically processes images uploaded to Amazon S3. When a user uploads an image to an S3 bucket, the system will trigger an AWS Lambda function that performs tasks such as resizing the image, compressing it, or applying filters.

You will use Amazon S3 to store both the original and processed images, and AWS Lambda will scale automatically based on the number of incoming image uploads. Additionally, you will integrate Amazon SNS to send notifications to the user once the image processing is complete.

This project will teach you how to leverage AWS's serverless capabilities to build an efficient, scalable image processing system.

SERVICES USED:

- ✓ Amazon S3
- ✓ AWS Lambda
- ✓ AWS IAM
- ✓ Amazon SNS(simple notification services)

DESCRIPTION:

Amazon S3: AWS S3 (Simple Storage Service) is a cloud data storage service. It is one of the most popular services of AWS. It has high scalability, availability, security and is cost effective. S3 has different storage tiers depending on the use case. Some common use cases of AWS S3 are:

1. **Storage:** It can be used for storing large amounts of data.
2. **Backup and Archive:** S3 has different storage tiers based on how frequent the data is accessed which can be used to backup critical data at low costs.
3. **Static website:** S3 offers static website hosting through HTML files stored in S3.
4. **Data lakes and big data analytics:** Companies can use [AWS S3](#) as a data lake and then run analytics on it for getting business insights and take critical decisions.

AWS Lambda:

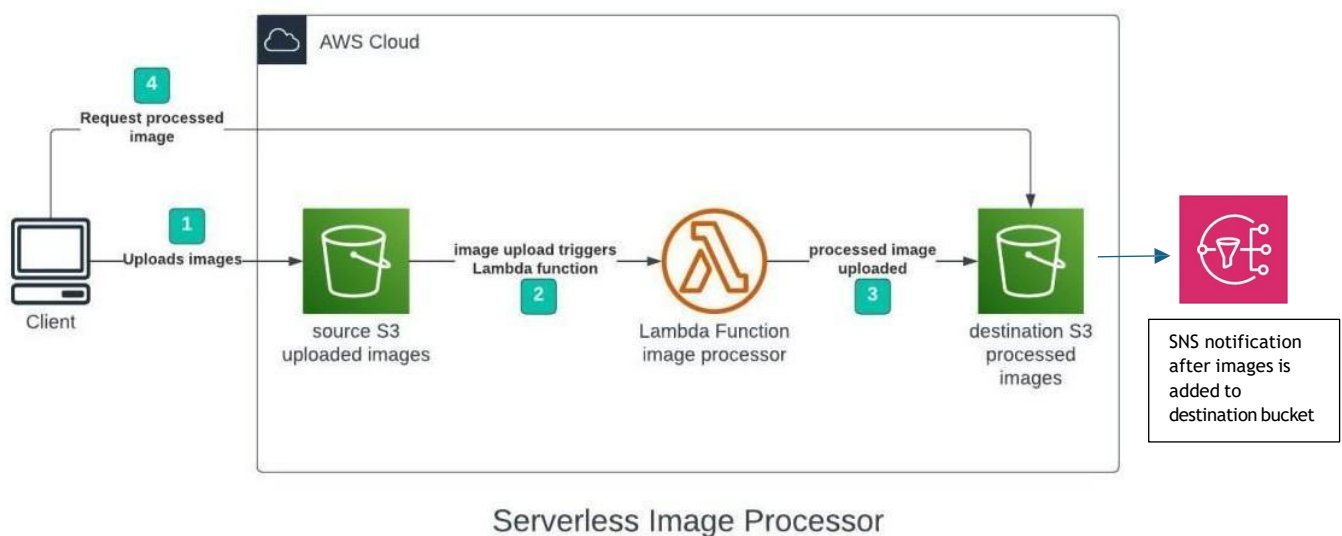
[AWS Lambda](#) is a serverless, event-driven compute service that lets you run code for virtually any type of application or backend service without provisioning or managing servers. Lambda functions run on demand i.e. they execute only when needed and you pay only for what you compute. Lambda is well integrated with many other AWS services. It supports a wide variety of programming languages.

Some common use cases for AWS Lambda are:

1. **File processing:** You can use Lambda for processing files as they are uploaded in an S3 bucket or whenever some event triggers the function.
2. **Data and analytics:** You can pass a data stream to your Lambda function and then create analysis from that.
3. **Website:** Lambda can also be used for creating websites. This is cost effective because you are charged only for the time when the servers are running.

Serverless Image Processing Flow

1. User uploads a file to the source S3 bucket (which is used for storing uploaded images).
2. When the image is uploaded to a source S3 bucket, it triggers an event which invokes the Lambda function. The lambda function processes the image.
3. Processed image is stored in the destination S3 bucket.
4. The processed image is requested by the user.



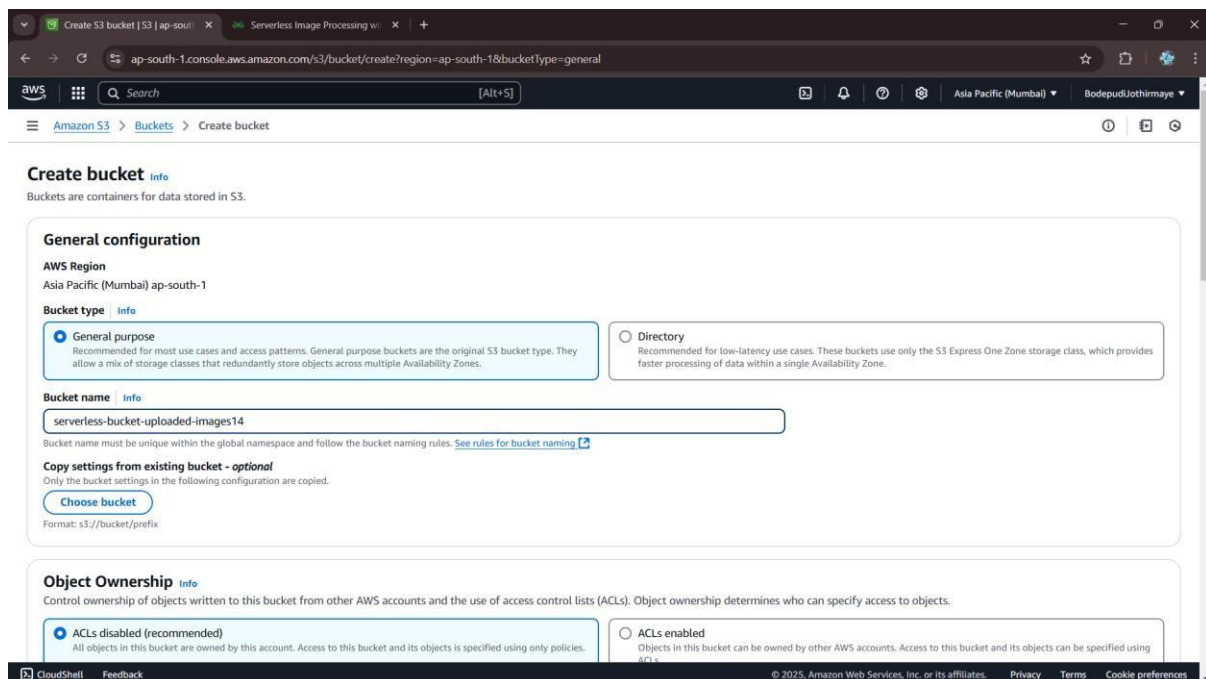
PROCEDURE:

Step 1 - Creating S3 buckets

We will use two S3 buckets:

1. **source Bucket:** For storing uploaded images.
2. **destination Bucket:** For storing processed images.

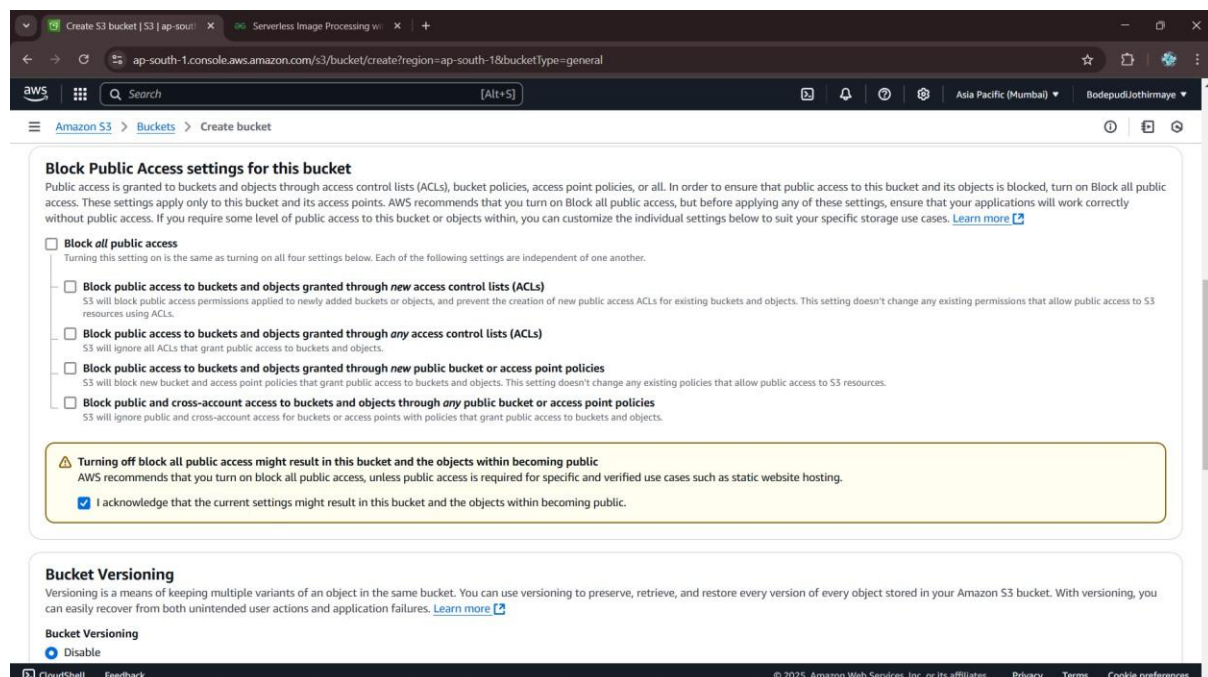
Go to S3 console and click Create bucket. Enter bucket name as 'serverless-bucket-uploaded- images'. Choose any AWS region as 'ap-south-1'.

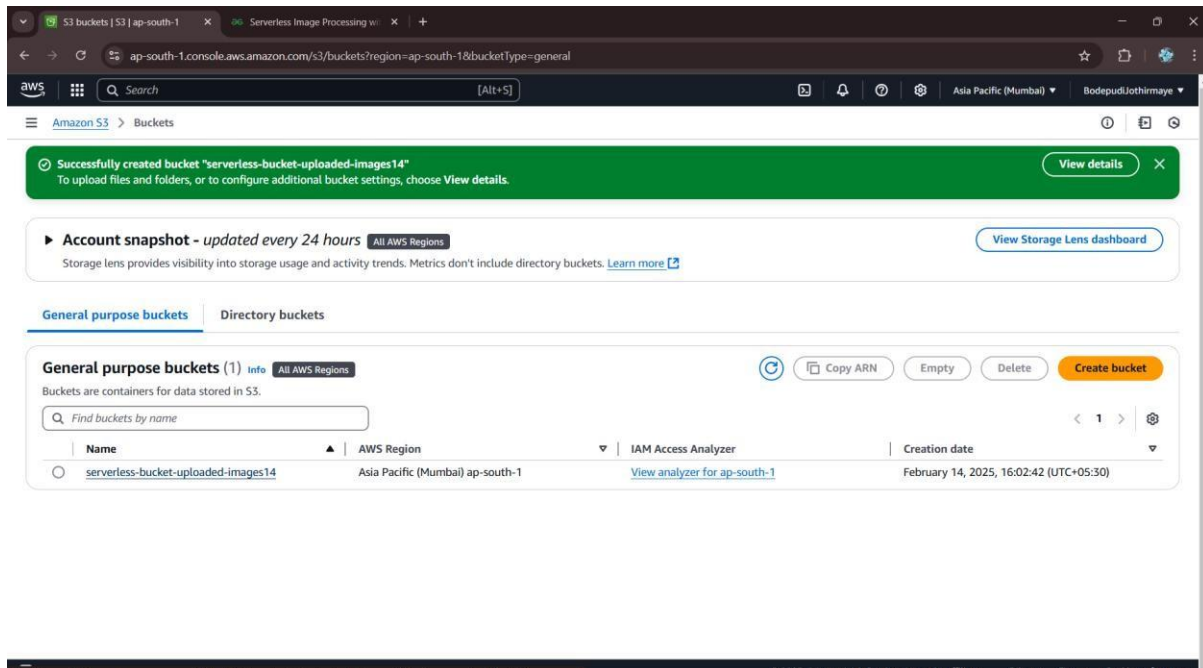


Step 2 - Configuring S3 bucket policy

In 'Block Public Access settings for this bucket' section disable "block all public access". You will get a warning that the bucket and its objects might become public. Agree to the warning. (**Note: we are making this bucket public only for this project, it is not recommended to make an S3 bucket public if not needed**).

Leave all other settings as default and create bucket. Similarly, create another bucket named 'serverless-bucket-processed-images' with the same region. This bucket will be used to store the processed images. Although we enabled public access while creating the buckets, we still need to attach a bucket policy to access the objects stored in it. (Policies in AWS are JSON documents which defines the permissions for performing actions on a certain resource.)





Go to your source bucket and then click on Permissions tab. Scroll to Bucket Policy tab. Click Edit. You will be redirected to the policy editor. Click on policy generator.

Enter the following settings:

- Type of policy: [S3 Bucket Policy](#)
- Effect: Allow
- Principal: *
- Actions: GetObject
- Amazon Resource Name (ARN):

arn:aws:s3:::SOURCE_BUCKET_NAME/*

SOURCE_BUCKET_NAME is the name of the bucket used for uploading the images.

Click Add Statement and then generate policy. Copy the JSON object.

Step 2: Add Statement(s)

A statement is the formal description of a single permission. See a description of elements that you can use in statements.

Effect ☒ Allow ☐ Deny

Principal

Use a comma to separate multiple values.

AWS Service ☐ All Services ("*")

Use multiple statements to add permissions for more than one service.

Actions ☐ All Actions ("*")

Amazon Resource Name (ARN)

ARN should follow the following format: arn:aws:s3:::{BucketName}/{KeyName}.
Use a comma to separate multiple values.

Add Conditions (Optional)

[Add Statement](#)

You added the following statements. Click the button below to Generate a policy.

Principal(s)	Effect	Action	Resource	Conditions
* *	Allow	s3:GetObject	arn:aws:s3:::serverless-bucket-uploaded-images14/*	None

Step 3: Generate Policy

A policy is a document (written in the Access Policy Language) that acts as a container for one or more statements.

[Generate Policy](#) [Start Over](#)

This AWS Policy Generator is provided for informational purposes only, you are still responsible for your use of Amazon Web Services technologies and ensuring that your use is in compliance with all applicable terms and conditions. This AWS Policy Generator is provided as is without warranty of any kind, whether express, implied, or statutory. This AWS Policy Generator does not create the legal basis for any Amazon Web Services transaction.

Paste it in the policy editor and then save changes.

Edit bucket policy - S3 bucket

ap-south-1.console.aws.amazon.com/s3/bucket/serverless-bucket-uploaded-images14/property/policy/edit?region=ap-south-1&bucketType=general

Amazon S3 > Buckets > serverless-bucket-uploaded-images14 > Edit bucket policy

Policy

```

1 {
2   "Id": "Policy1739529296274",
3   "Version": "2012-10-17",
4   "Statement": [
5     {
6       "Sid": "Stmt1739529267133",
7       "Action": [
8         "s3:GetObject"
9       ],
10      "Effect": "Allow",
11      "Resource": "arn:aws:s3:::serverless-bucket-uploaded-images14/*",
12      "Principal": "*"
13    }
14  ]
15 }
```

[+ Add new statement](#)

Edit statement

Select a statement

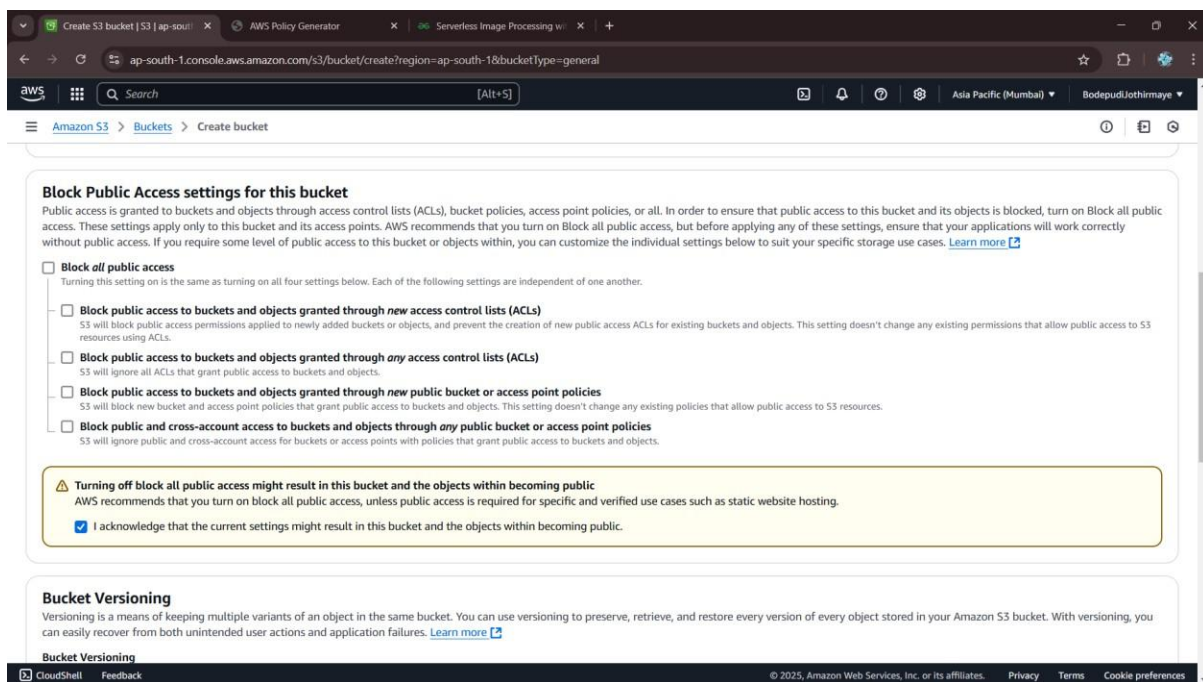
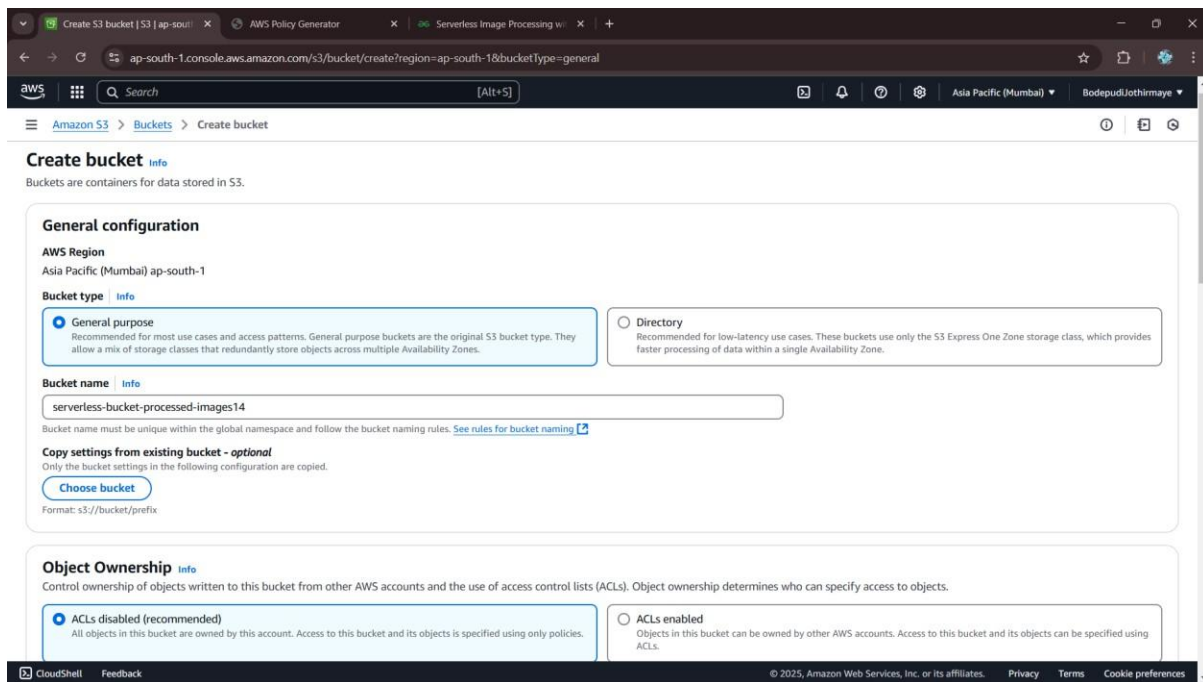
Select an existing statement in the policy or add a new statement.

[+ Add new statement](#)

[+ Add new statement](#)

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Repeat the above steps and create other Bucket named serverless-bucket-processed-images



Paste it in the policy editor and then save changes.

Follow same steps to attach a policy to the processed images S3 bucket. The policy settings for destination bucket are:

- Type of policy: S3 Bucket Policy
- Effect: Allow

- Principal:*
- Actions: GetObject, PutObject, and PutObjectAcl
- Amazon Resource Name (ARN): arn:aws:s3:::DESTINATION_BUCKET_NAME/*

DESTINATION_BUCKET_NAME is the name of the bucket used for storing processed images.

The screenshot shows the AWS Policy Generator interface. The 'AWS Service' is set to 'Amazon S3'. The 'Actions' are 's3:GetObject', 's3:PutObject', and 's3:PutObjectAcl'. The 'Amazon Resource Name (ARN)' is 'arn:aws:s3:::serverless-bucket-processed-images14/*'. The 'Conditions' are set to 'None'. The 'Principal(s)' is '*'. The 'Effect' is 'Allow'. The 'Generate Policy' button is highlighted.

Principal(s)	Effect	Action	Resource	Conditions
*	Allow	s3:GetObject s3:PutObject s3:PutObjectAcl	arn:aws:s3:::serverless-bucket-processed-images14/*	None

Step 3: Generate Policy

A policy is a document (written in the Access Policy Language) that acts as a container for one or more statements.

[Generate Policy](#) [Start Over](#)

This AWS Policy Generator is provided for informational purposes only, you are still responsible for your use of Amazon Web Services technologies and ensuring that your use is in compliance with all applicable terms and conditions. This AWS Policy Generator is provided as is without warranty of any kind, whether express, implied, or statutory. This AWS Policy Generator does not modify the applicable terms and conditions governing your use of Amazon Web Services technologies.

©2010, Amazon Web Services LLC or its affiliates. All rights reserved.

An amazon.com company

The screenshot shows the AWS console 'Bucket policy' page for the bucket 'serverless-bucket-processed-images14'. The policy is displayed in JSON format. The 'Edit' and 'Delete' buttons are visible. The 'Copy' button is also present.

```

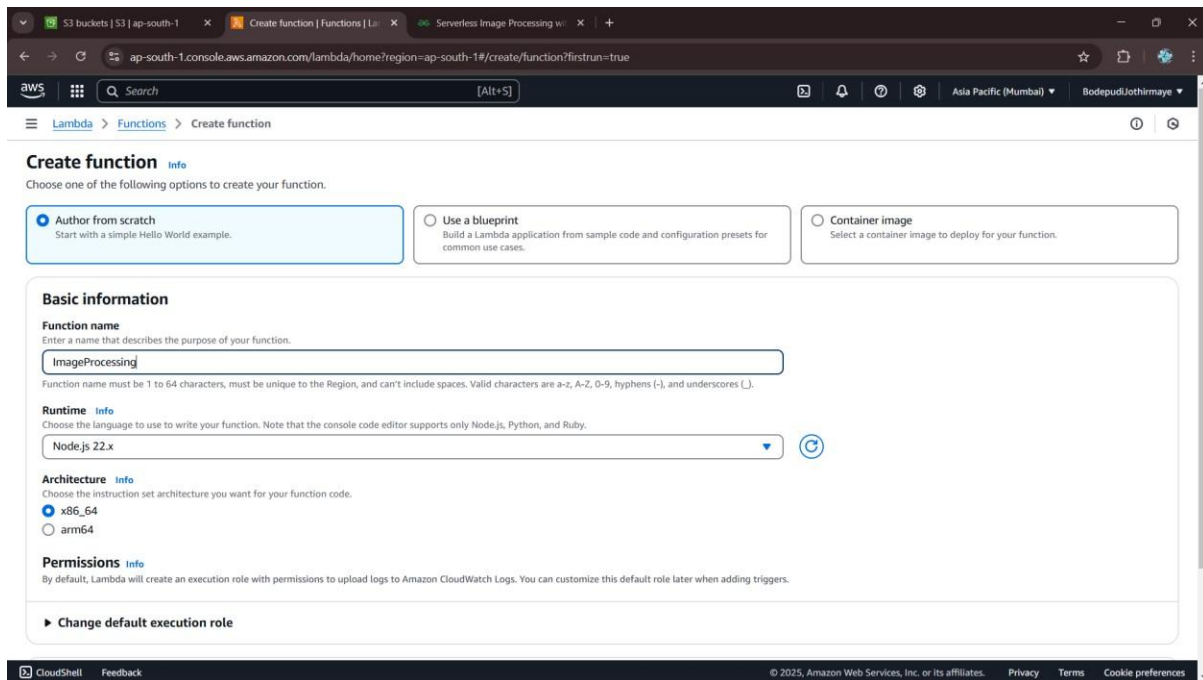
{
  "Version": "2012-10-17",
  "Id": "Policy1739529518043",
  "Statement": [
    {
      "Sid": "Stmt1739529503141",
      "Effect": "Allow",
      "Principal": "*",
      "Action": [
        "s3:GetObject",
        "s3:PutObject",
        "s3:PutObjectAcl"
      ],
      "Resource": "arn:aws:s3:::serverless-bucket-processed-images14/*"
    }
  ]
}

```

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 3 - Creating Lambda function

Go to AWS Lambda console. Navigate to Functions section. Click Create Function and name it "ImageProcessing". Select runtime as "NodeJS 22.x" and architecture as "x86_64". Leave all other settings as default. Create the function.



The screenshot shows the AWS Lambda 'Create function' page. The 'Author from scratch' option is selected. The function name is 'ImageProcessing', the runtime is 'Node.js 22.x', and the architecture is 'x86_64'. The page includes a 'Basic information' section with fields for function name, runtime, and architecture. A link to 'Change default execution role' is visible at the bottom of the form.

In the code editor on the Lambda function page paste the following code. This function is executed whenever an image is uploaded to our source S3 bucket and creates two images (thumbnail (300x300) and coverphoto(800x800)) and stores it in the destination S3 bucket. **(Note: The value of processedImageBucket in the code should be set to the name of the destination bucket).**

CODE:

```
import sharp from "sharp";
import path from "path";
import { S3 } from "@aws-sdk/client-s3";
import { SNS } from "@aws-sdk/client-sns";
// Set up AWS clients
const s3 = new S3({ region: "ap-south-1" });
const sns = new SNS({ region: "ap-south-1" });
```

```

// Configuration

const processedImageBucket = "serverless-bucket-processed-images14";

const snsTopicArn = "arn:aws:sns:ap-south-1:235494801123:ImageProcessorNotification"; //
Replace with your SNS Topic ARN


export const handler = async (event) => {
  console.log("Received S3 Event:", JSON.stringify(event));

  for (const record of event.Records) {
    console.log("Processing Record:", record);

    const bucketName = record.s3.bucket.name;
    const objectKey = decodeURIComponent(record.s3.object.key.replace(/\+/g, " "));
    const { name: fileName, ext: fileExt } = path.parse(objectKey);

    try {
      // Fetch the image from S3
      const imageObject = await s3.getObject({ Bucket: bucketName, Key: objectKey });
      const imageBuffer = await streamToBuffer(imageObject.Body);

      // Generate processed images
      const resizedThumbnail = await resizeImage(imageBuffer, 300, 300);
      const resizedCoverPhoto = await resizeImage(imageBuffer, 800, 800);

      console.log("Images resized successfully");

      // Upload processed images to S3
      await uploadToS3(`${fileName}_thumbnail${fileExt}`, resizedThumbnail);
      await uploadToS3(`${fileName}_coverphoto${fileExt}`, resizedCoverPhoto);

      // Send SNS Notification

```

```

        await sendSNSNotification(`Image processing completed for ${objectKey}`);

    } catch (error) {
        console.error("Error processing image:", error);
        await sendSNSNotification(`Image processing failed for ${objectKey}: ${error.message}`);
    }
}
};

```

// Helper: Convert readable stream to buffer

```

const streamToBuffer = async (stream) => {
    const chunks = [];
    for await (const chunk of stream) {
        chunks.push(chunk);
    }
    return Buffer.concat(chunks);
};

```

// Helper: Resize image using Sharp

```

const resizeImage = async (imageBuffer, width, height) => {
    return await sharp(imageBuffer)
        .resize({ width, height, fit: sharp.fit.cover })
        .withMetadata()
        .toBuffer();
};

```

// Helper: Upload image to S3

```

const uploadToS3 = async (fileName, imageBuffer) => {
    await s3.putObject({
        Bucket: processedImageBucket,
        Key: fileName,

```



```

    Body: imageBuffer,

    CacheControl: "max-age=3600",

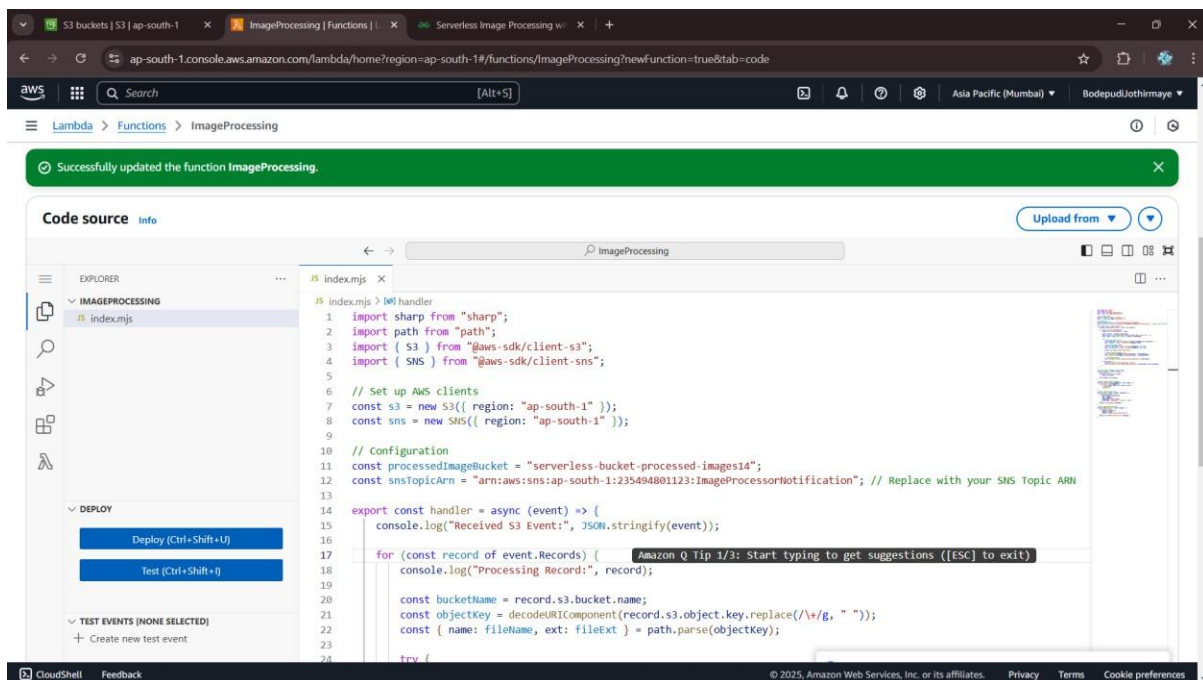
    ContentType: "image/jpeg", // Modify if needed
  });

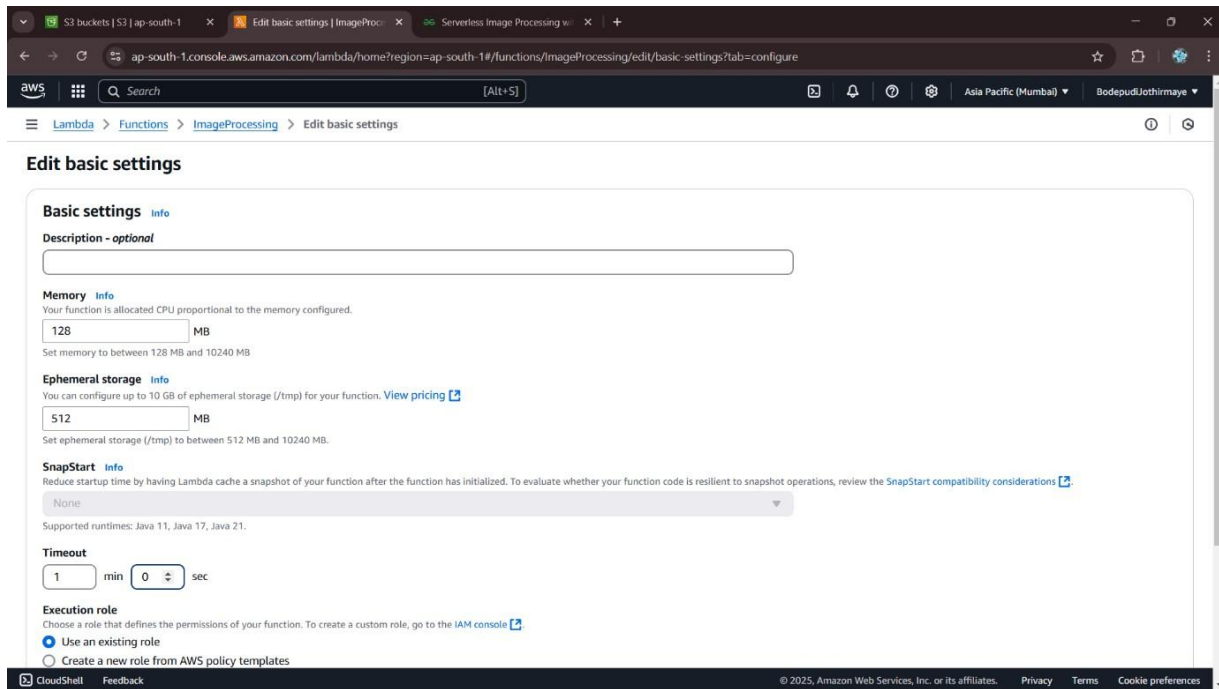
  console.log(`Uploaded: ${fileName}`);
};

// Helper: Send SNS Notification
const sendSNSNotification = async (message) => {
  await sns.publish({
    TopicArn: snsTopicArn,
    Message: message,
    Subject: "Image Processing Notification",
  });
  console.log(`SNS Notification Sent: ${message}`);
};

Save and click deploy changes.

```





Go to Configuration tab and Edit the general configuration. There set the timeout to 1 min (timeout is the maximum time for which a Lambda function will run after which it stops running). We need to increase the timeout because the image can take time to process. Click on Save changes.

Step 4 - Creating Lambda layer and attaching it to Lambda function

Layers in Lambda is used to add dependencies to a Lambda Function. Lambda Layers reduces the code size of Lambda functions as we do not need to upload the dependencies with the function. It also useful for code reusability as we can reuse the layer with multiple functions if they require the same dependencies.

First we need to create a zip file with all the dependencies (node modules in our case) required by our Lambda function.

Create a folder "aws-serverless-image-processor". Inside this directory create another directory "nodejs" (it is compulsory to name this as "nodejs"). Open a terminal and go to nodejs directory. Install *sharp* module with the following command (platform is linux because the Lambda function runs on a Linux machine so we require the node_modules for Linux).

```
npm install --arch=x64 --platform=linux sharp
```

Now create a zip file of the nodejs directory and name it "sharp-layer.zip".

Go to Layers in Lambda console. Click Create layer. Name it "sharp-layer".

Upload your nodejs "sharp-layer.zip" file here. Select x86_64 architecture. Select NodeJS 22.x in compatible runtimes. Click on Create Layer.

Layer configuration

Name
sharp-layer

Description - optional

☒ Upload a .zip file
☐ Upload a file from Amazon S3

Choose file
sharp-layer.zip
7.45 MB

For files larger than 10 MB, consider uploading using Amazon S3.

Compatible architectures - optional | Info
Choose the compatible instruction set architectures for your layer.
x86_64

Compatible runtimes - optional | Info
Choose up to 15 runtimes.
Node.js 22.x

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

sharp-layer Delete Download Create version

Version details

Version 5	Version ARN arn:aws:lambda:ap-south-1:235494801123:layers:sharp-layer:5	Description -
Created 25 seconds ago	License -	Compatible runtimes nodejs22.x
Compatible architectures x86_64		

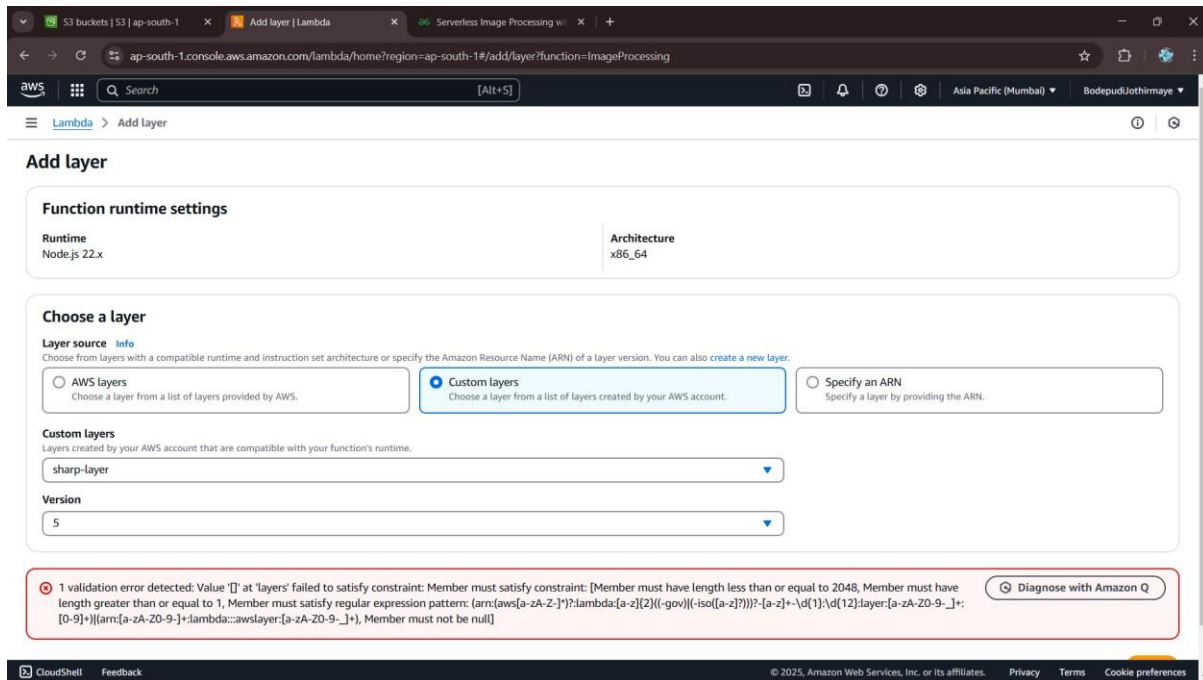
Versions Functions using this version

All versions (1) Last fetched 21 seconds ago

Filter by attributes or search by keyword

Version	Version ARN	Description
5	arn:aws:lambda:ap-south-1:235494801123:layers:sharp-layer:5	-

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

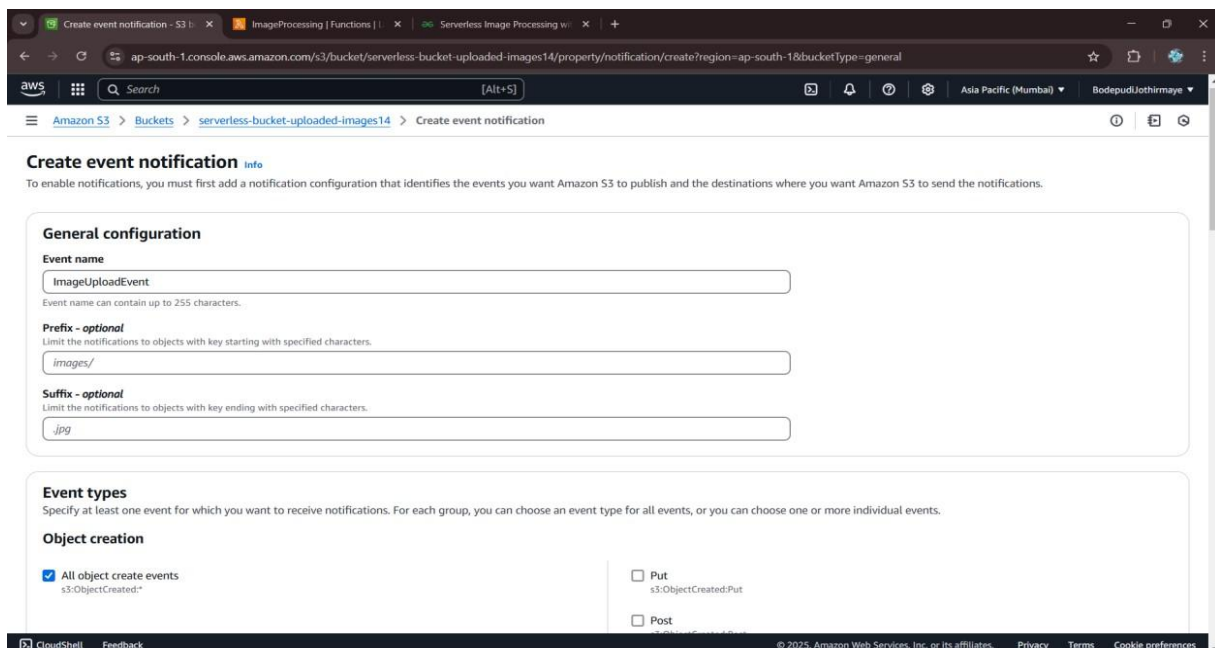


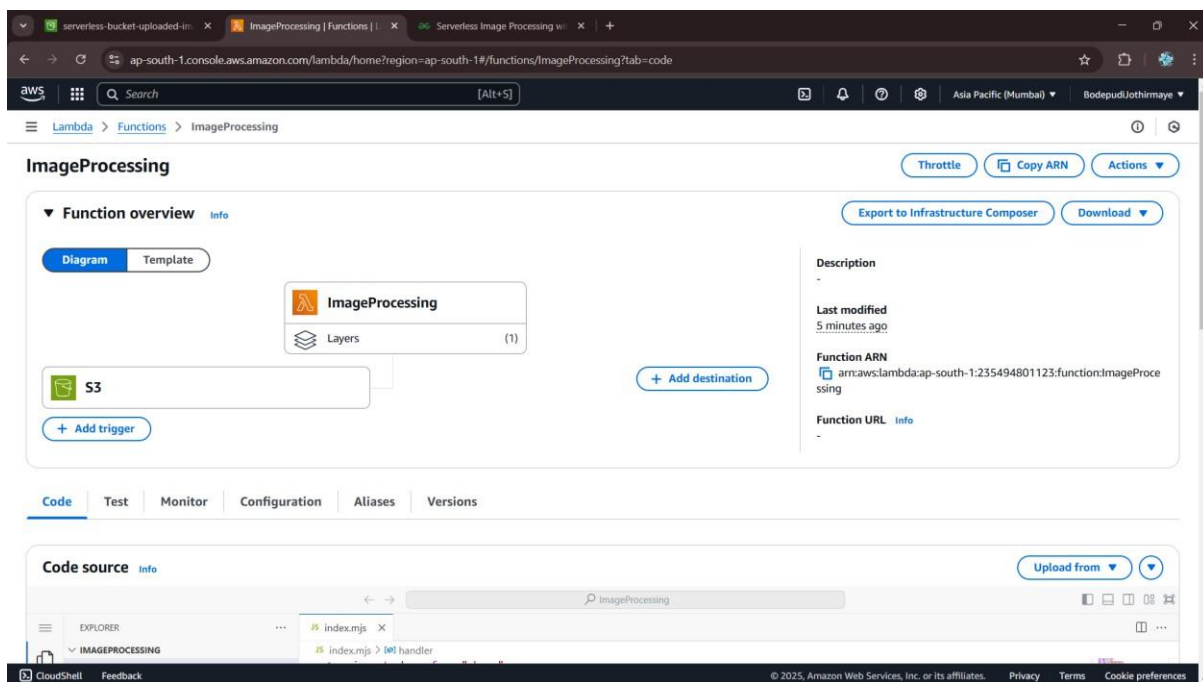
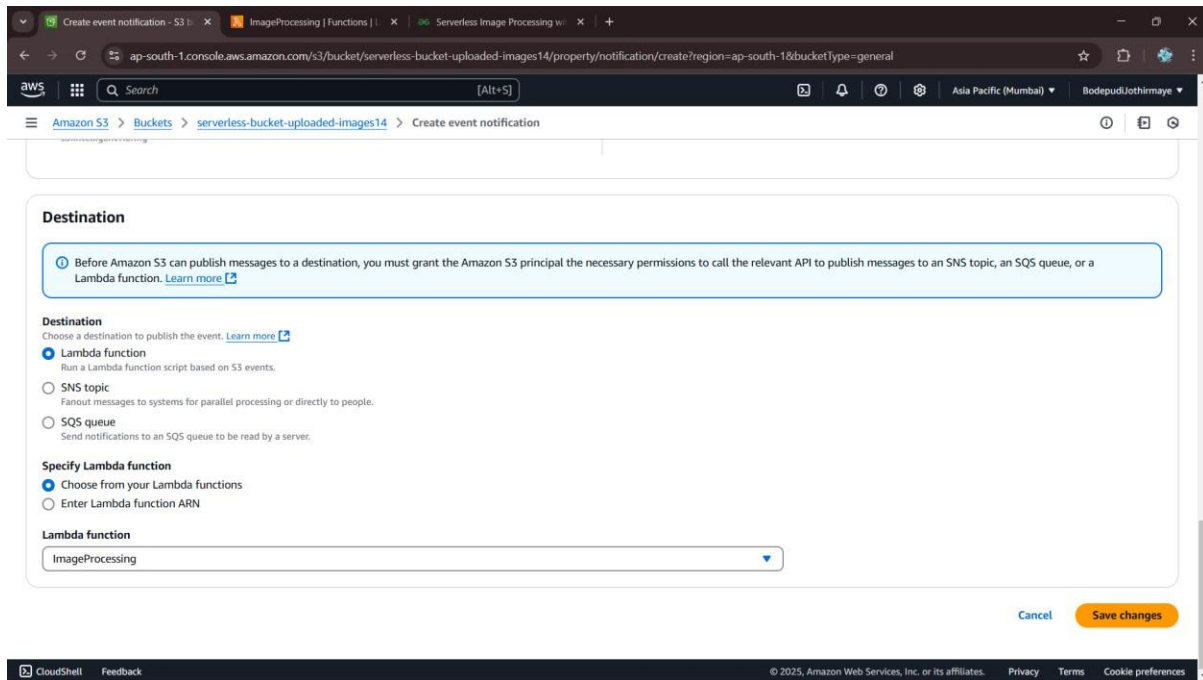
Step 5 - Creating S3 trigger

Now we need our Lambda function to know when an image is uploaded to the source S3 bucket. We can do this by adding an event to the source S3 bucket and configure it to get triggered when an image is uploaded to the bucket which in turn invokes the Lambda function.

Go to S3 console. Select the source bucket ("serverless-bucket-uploaded-images"). Go to the Properties tab. Navigate to "Event Notifications". Click "Create Event Notifications".

Give an appropriate name to the event. Check the "All object create events".





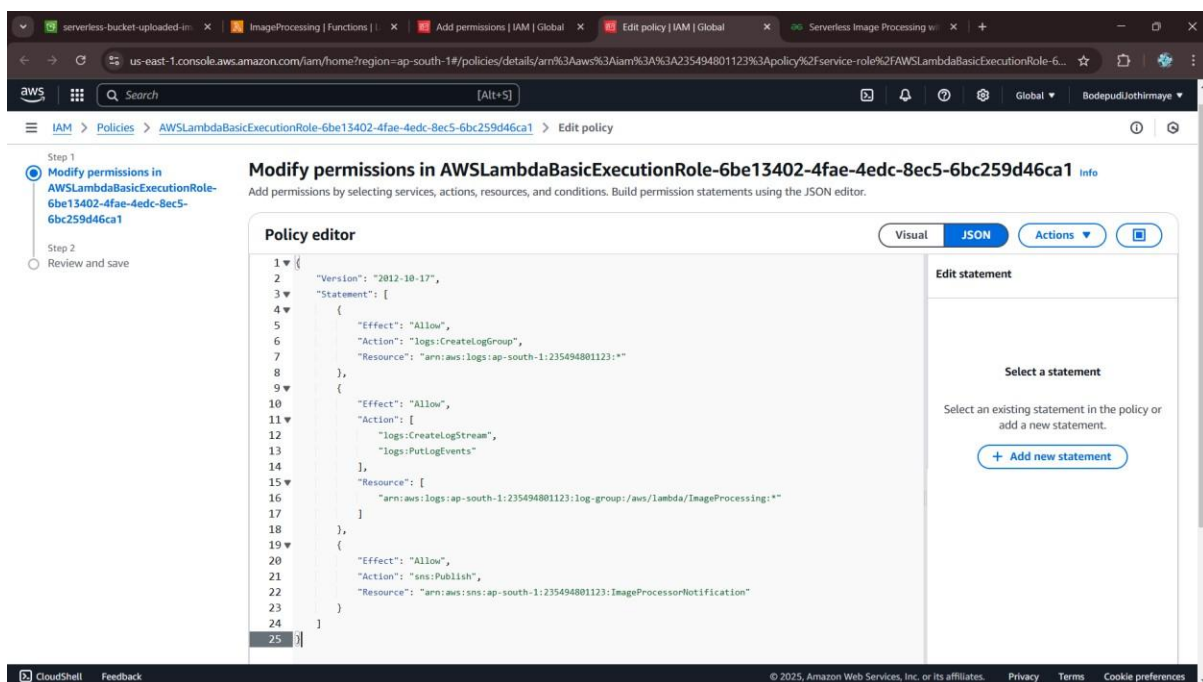
STEP 6: In IAM go to roles and click on “Attach Policy” and click on edit policy and add the following

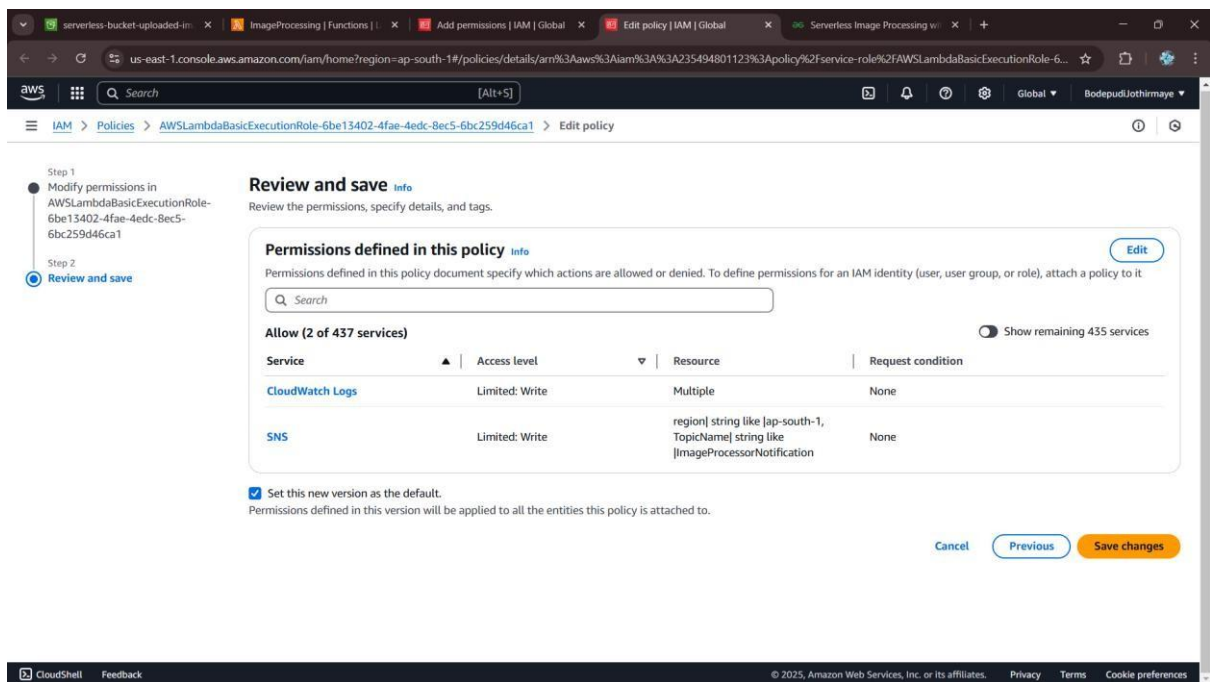
```
{
  "Version": "2012-10-17",
  "Statement": [
    {
```

```

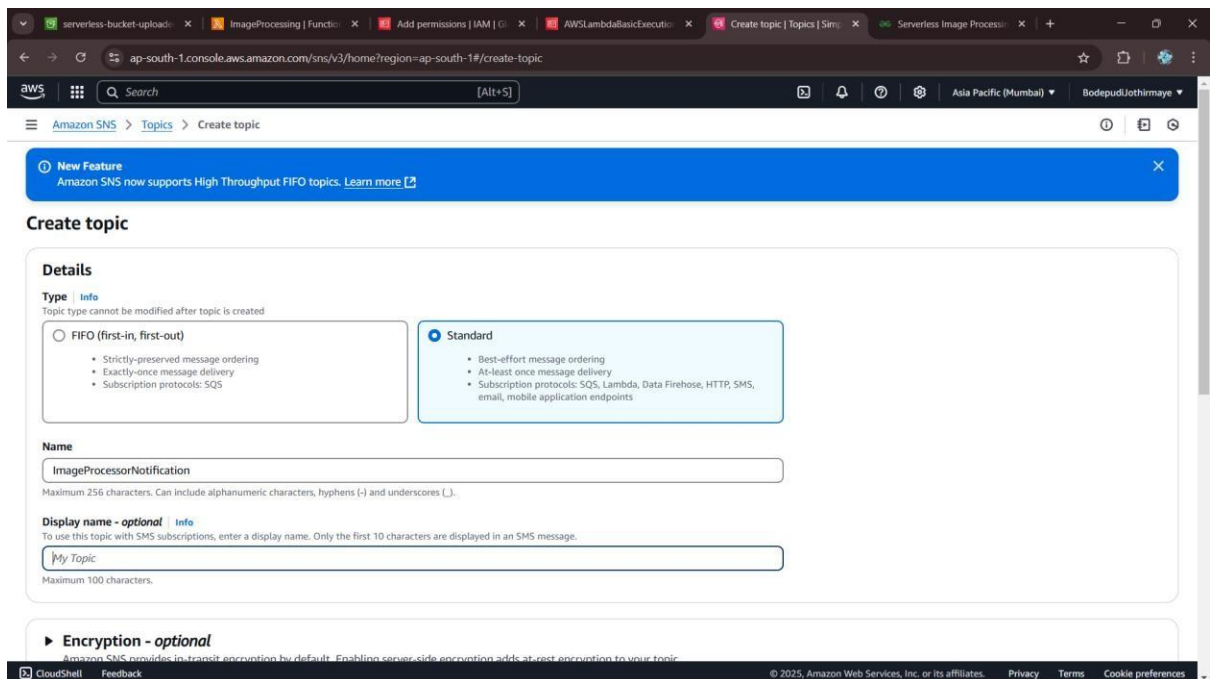
    "Effect": "Allow",
    "Action": "logs:CreateLogGroup",
    "Resource": "arn:aws:logs:ap-south-1:235494801123:*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "logs:CreateLogStream",
      "logs:PutLogEvents"
    ],
    "Resource": [
      "arn:aws:logs:ap-south-1:235494801123:log-group:/aws/lambda/ImageProcessing:*"
    ]
  },
  {
    "Effect": "Allow",
    "Action": "sns:Publish",
    "Resource": "arn:aws:sns:ap-south-1:235494801123:ImageProcessorNotification"
  }
]
}

```

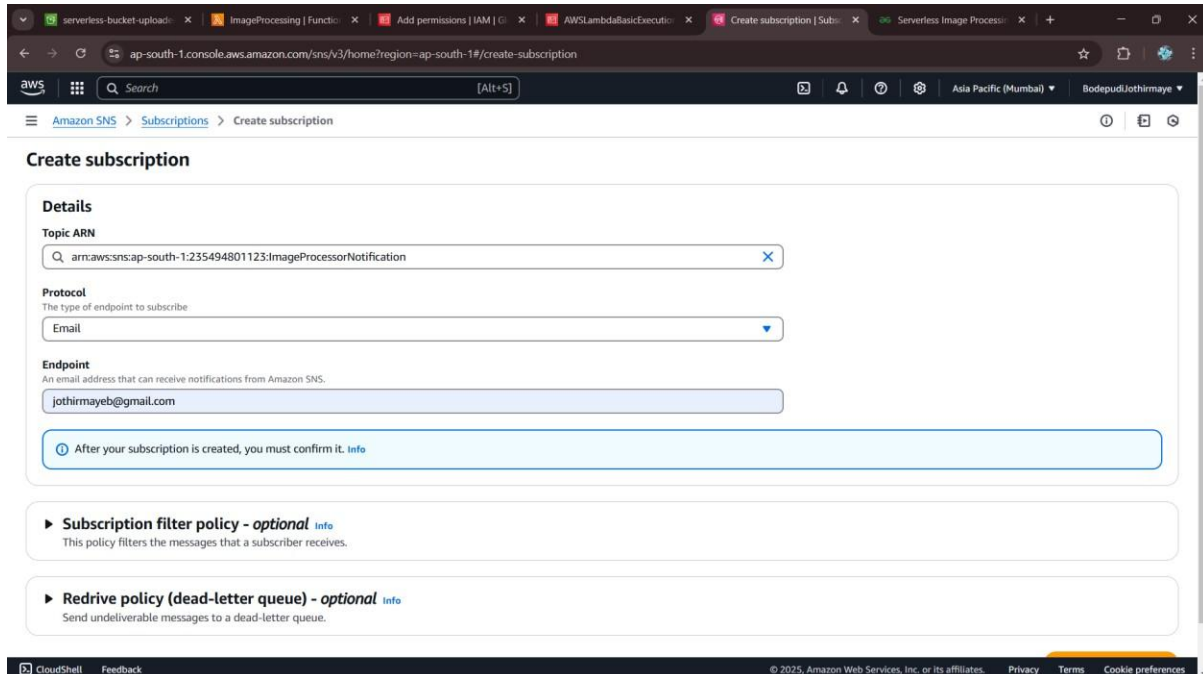




STEP 7: Create SNS topic named “ImageProcessorNotification” and click on save.

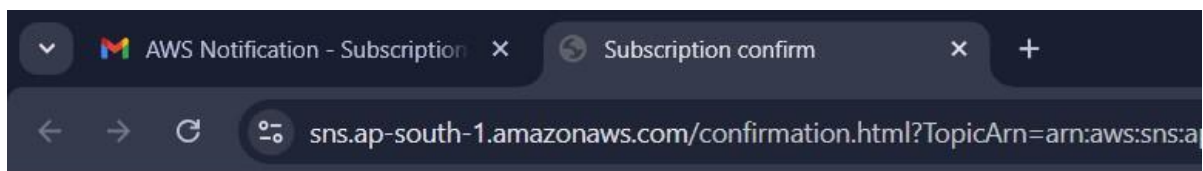


Now click on create new subscription and select protocol as email and add you email and click on create new subscription



The screenshot shows the AWS console 'Create subscription' page. The browser tabs include 'serverless-bucket-upload', 'ImageProcessing [Func...', 'Add permissions [IAM | G...', 'AWSLambdaBasicExecuti...', 'Create subscription [Sub...', and 'Serverless Image Process...'. The URL is 'ap-south-1.console.aws.amazon.com/sns/v3/home?region=ap-south-1#/create-subscription'. The page title is 'Create subscription'. The 'Details' section has a 'Topic ARN' field with the value 'arn:aws:sns:ap-south-1:235494801123:ImageProcessorNotification'. The 'Protocol' dropdown is set to 'Email'. The 'Endpoint' field contains 'jothirmayeb@gmail.com'. A note states: 'After your subscription is created, you must confirm it.' Below this are two optional sections: 'Subscription filter policy - optional' and 'Redrive policy (dead-letter queue) - optional'.

An email is sent to you to confirm the subscription



Simple Notification Service

Subscription confirmed!

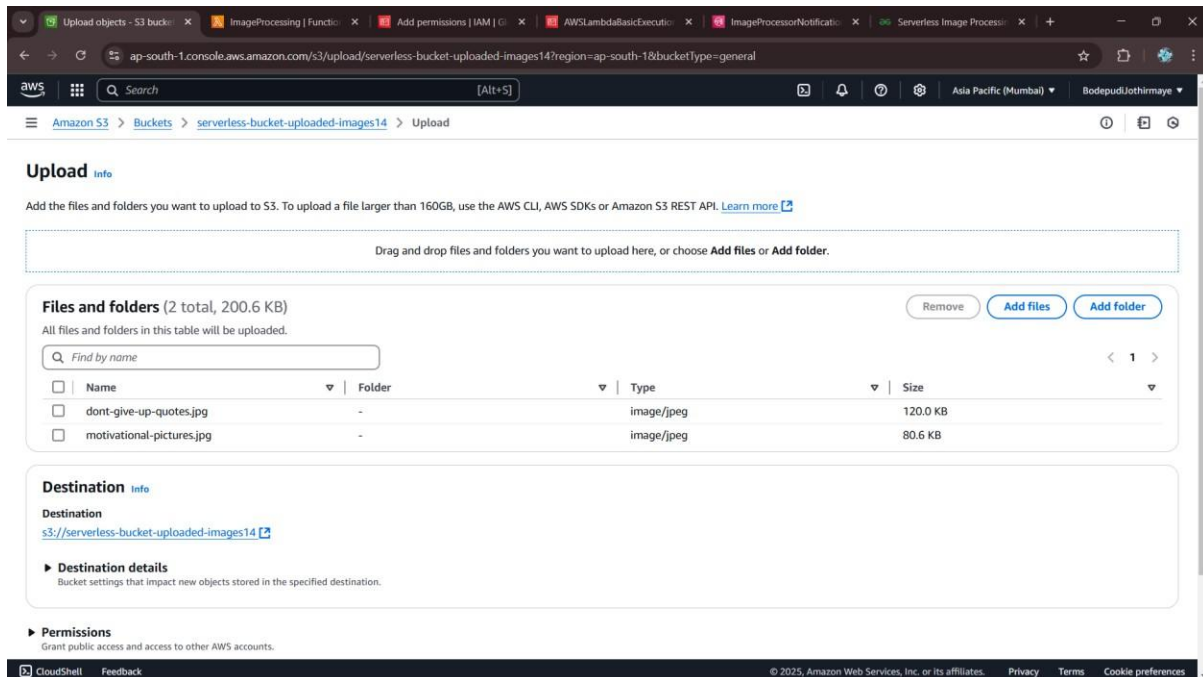
You have successfully subscribed.

Your subscription's id is:

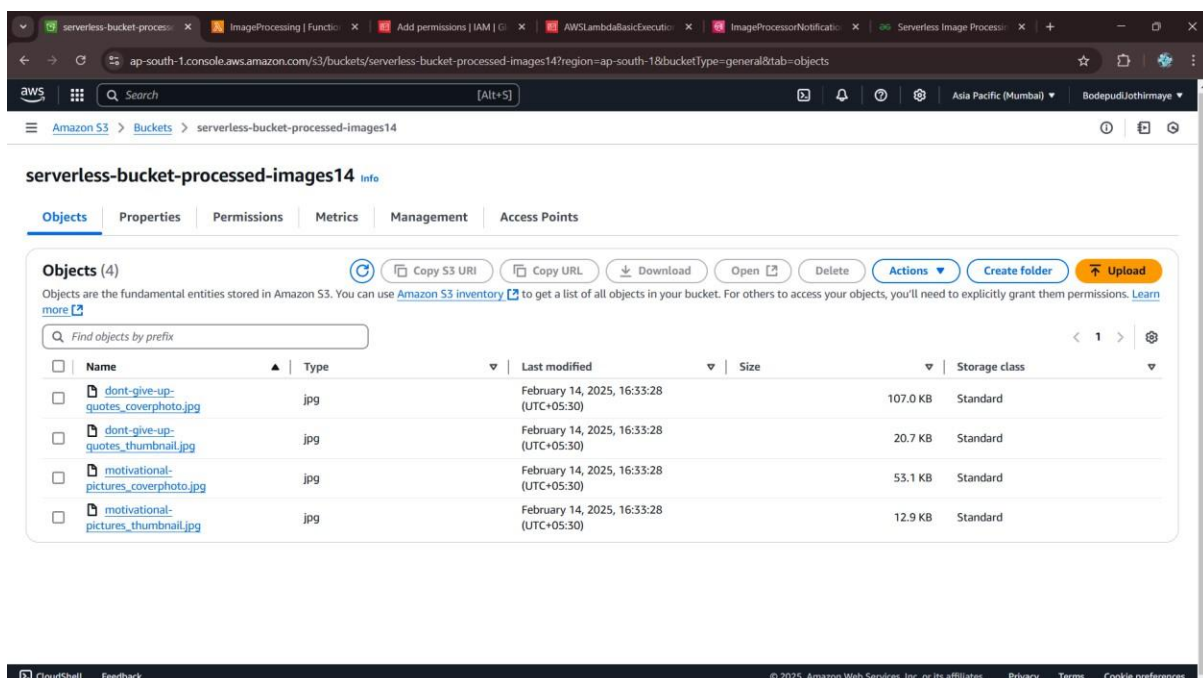
arn:aws:sns:ap-south-1:235494801123:ImageProcessorNotification:1b664540-446d-46f8-8aea-d44790fb645b

If it was not your intention to subscribe, [click here to unsubscribe](#).

Now add the images to serverless-bucket-upload-images14

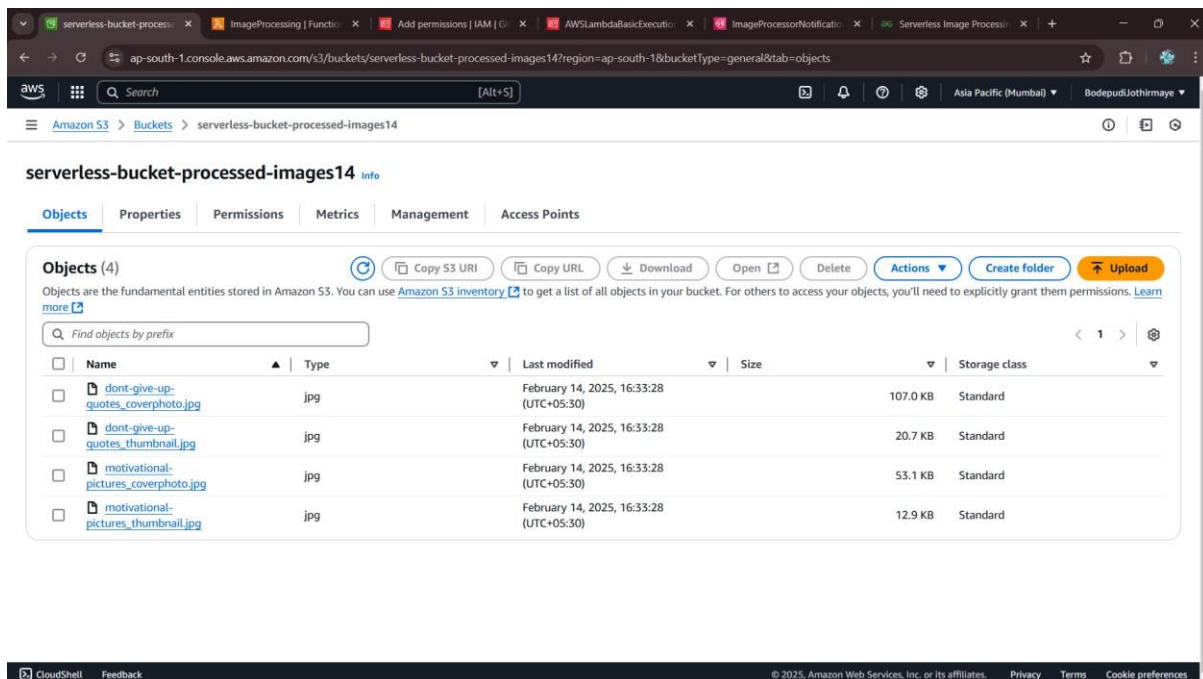
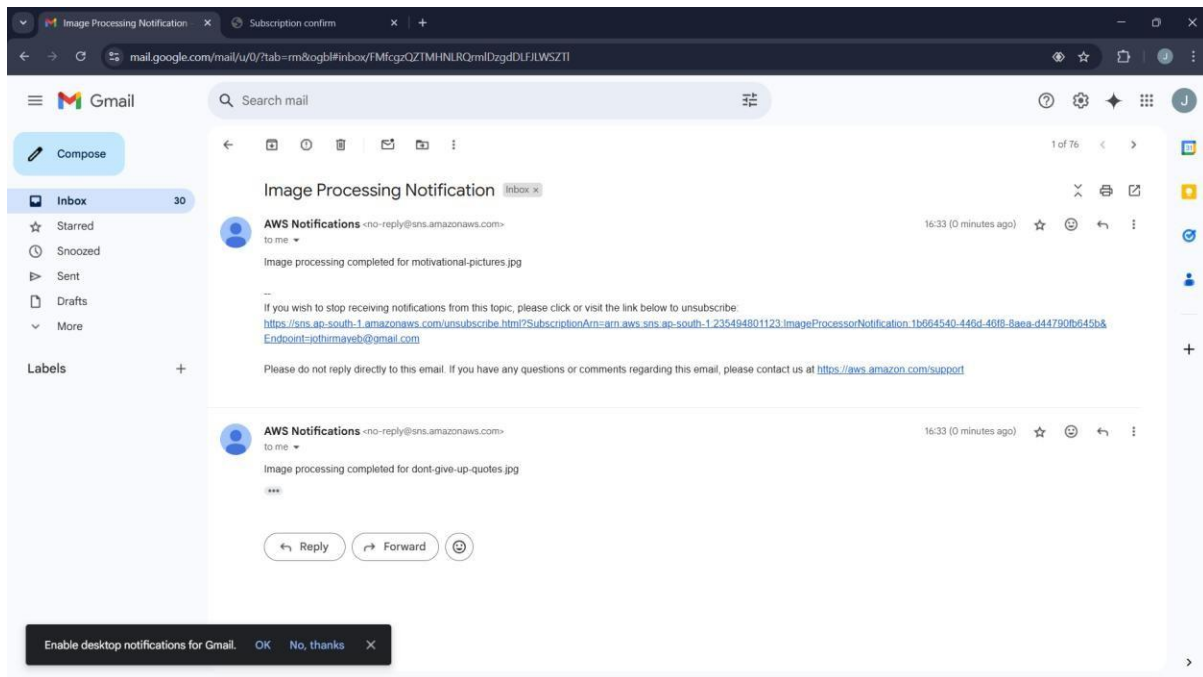


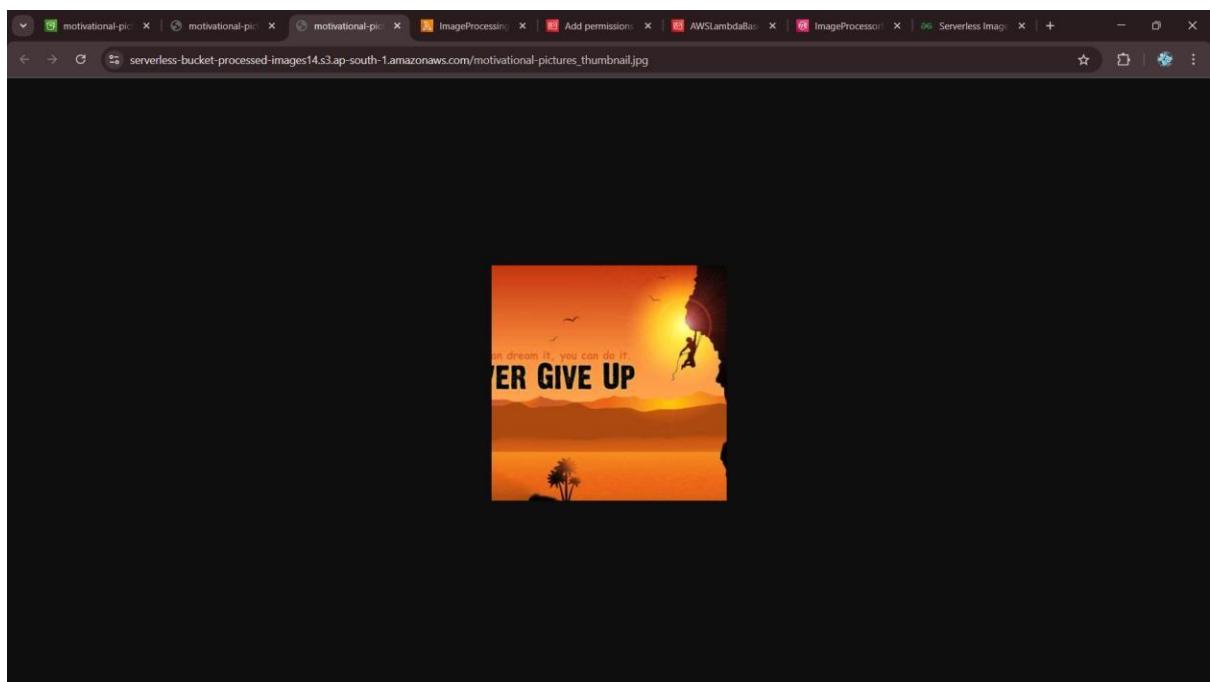
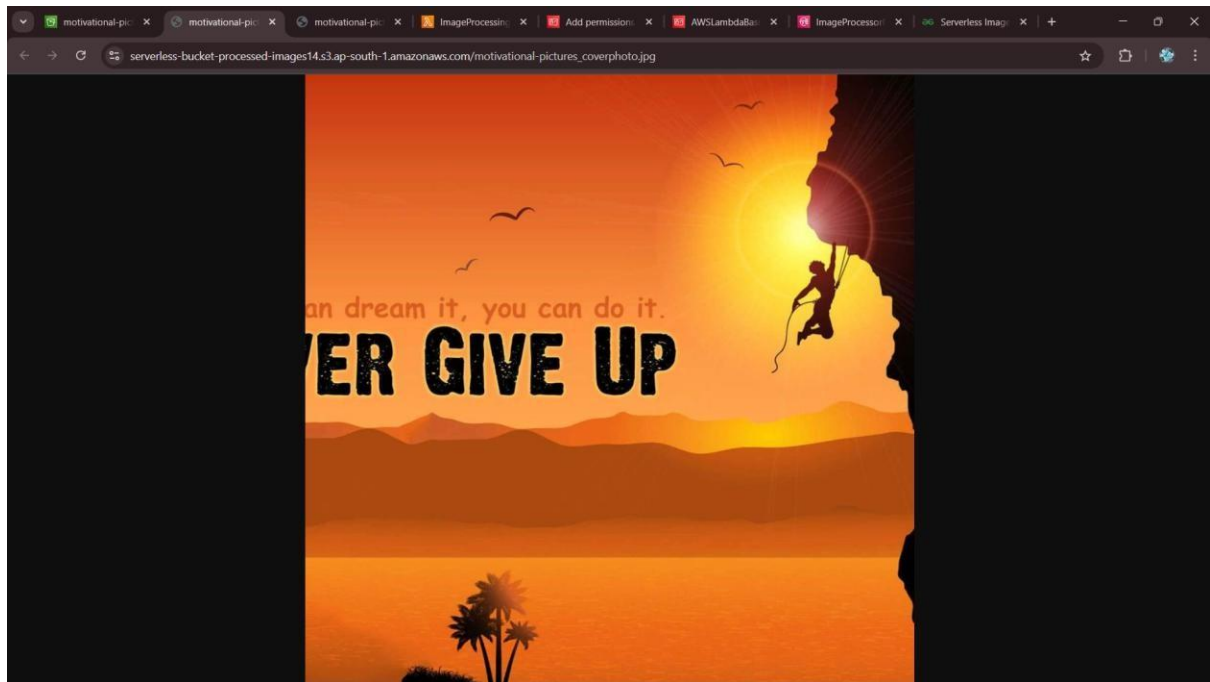
After the successful process of you image the two different sizes of the images will be added to your destination bucket.



OUTPUT:

After the successful process an mail will be sent to your email saying “image processing completed for image.jpg”





CONCLUSION:

Serverless image processing with AWS Lambda and S3 provides a scalable, cost-effective, and efficient solution for handling image transformations in the cloud. By leveraging AWS Lambda, you can automatically process images in response to S3 events without managing servers. Integration with other AWS services, such as Amazon S3, Amazon SNS, and Amazon CloudWatch, enables automated workflows, monitoring, and alerting.

This approach is ideal for real-time applications like thumbnail generation, watermarking, and format conversions. It reduces infrastructure costs, enhances scalability, and simplifies deployment. However, developers must optimize Lambda execution time, handle cold starts efficiently, and manage permissions securely.

By adopting serverless architecture for image processing, businesses can achieve high availability, fault tolerance, and improved performance, making it a practical choice for modern cloud-based applications.